

Efficient Data Import in R

Michael Otterstatter
BCCDC Biostats Session
Feb 15, 2019

Outline

- Data import tools in tidyverse
- Reading csv files
- Formatting columns
- Dealing with missing values
- Dealing with dates
- Reading other file types

Importing data

- Regardless of the analytic task, we almost always begin by taking data from an existing source (file, webpage, database) and importing it into an object in R, SAS, etc. for further manipulation and analysis

Importing data with `readr`

- `readr` can import from 7 different file formats:
 - `read_csv()`: comma separated (csv) files
 - `read_tsv()`: tab separated files
 - `read_delim()`: general delimited files
 - `read_fwf()`: fixed width files
 - `read_table()`: tabular files with columns separated by white-space.
 - `read_log()`: web log files

All have similar form and syntax, so we will focus on `read_csv` the most commonly used example

Data Import :: CHEAT SHEET



R's **tidyverse** is built around **tidy data** stored in **tibbles**, which are enhanced data frames.



The front side of this sheet shows how to read text files into R with **readr**.



The reverse side shows how to create tibbles with **tibble** and to layout tidy data with **tidyr**.

OTHER TYPES OF DATA

Try one of the following packages to import other types of files

- **haven** - SPSS, Stata, and SAS files
- **readxl** - excel files (.xls and .xlsx)
- **DBI** - databases
- **jsonlite** - json
- **xml2** - XML
- **httr** - Web APIs
- **rvest** - HTML (Web Scraping)

Save Data

Save **x**, an R object, to **path**, a file path, as:

Comma delimited file

```
write_csv(x, path, na = "NA", append = FALSE,
  col_names = !append)
```

File with arbitrary delimiter

```
write_delim(x, path, delim = " ", na = "NA",
  append = FALSE, col_names = !append)
```

CSV for excel

```
write_excel_csv(x, path, na = "NA", append =
  FALSE, col_names = !append)
```

String to file

```
write_file(x, path, append = FALSE)
```

String vector to file, one element per line

```
write_lines(x, path, na = "NA", append = FALSE)
```

Object to RDS file

```
write_rds(x, path, compress = c("none", "gz",
  "bz2", "xz"), ...)
```

Tab delimited files

```
write_tsv(x, path, na = "NA", append = FALSE,
  col_names = !append)
```

Read Tabular Data - These functions share the common arguments:

```
read_(file, col_names = TRUE, col_types = NULL, locale = default_locale(), na = c("", "NA"),
  quoted_na = TRUE, comment = "", trim_ws = TRUE, skip = 0, n_max = Inf, guess_max = min(1000,
  n_max), progress = interactive())
```

a,b,c
1,2,3
4,5,NA

A	B	C
1	2	3
4	5	NA

Comma Delimited Files

read_csv("file.csv")

To make file.csv run:

```
write_file(x = "a,b,c\n1,2,3\n4,5,NA", path = "file.csv")
```

a;b;c
1;2;3
4;5;NA

A	B	C
1	2	3
4	5	NA

Semi-colon Delimited Files

read_csv2("file2.csv")

```
write_file(x = "a;b;c\n1;2;3\n4;5;NA", path = "file2.csv")
```

a|b|c
1|2|3
4|5|NA

A	B	C
1	2	3
4	5	NA

Files with Any Delimiter

read_delim("file.txt", delim = "|")

```
write_file(x = "a|b|c\n1|2|3\n4|5|NA", path = "file.txt")
```

a b c
1 2 3
4 5 NA

A	B	C
1	2	3
4	5	NA

Fixed Width Files

read_fwf("file.fwf", col_positions = c(1, 3, 5))

```
write_file(x = "a b c\n1 2 3\n4 5 NA", path = "file.fwf")
```

Tab Delimited Files

read_tsv("file.tsv") Also **read_table()**.

```
write_file(x = "a\tb\tc\n1\t2\t3\n4\t5\tNA", path = "file.tsv")
```

USEFUL ARGUMENTS

a,b,c
1,2,3
4,5,NA

Example file

```
write_file("a,b,c\n1,2,3\n4,5,NA","file.csv")  
f <- "file.csv"
```

1	2	3
4	5	NA

Skip lines

```
read_csv(f, skip = 1)
```

A	B	C
1	2	3
4	5	NA

No header

```
read_csv(f, col_names = FALSE)
```

A	B	C
1	2	3

Read in a subset

```
read_csv(f, n_max = 1)
```

x	y	z
A	B	C
1	2	3
4	5	NA

Provide header

```
read_csv(f, col_names = c("x", "y", "z"))
```

A	B	C
NA	2	3
4	5	NA

Missing Values

```
read_csv(f, na = c("1", ""))
```

Read Non-Tabular Data

Read a file into a single string

```
read_file(file, locale = default_locale())
```

Read each line into its own string

```
read_lines(file, skip = 0, n_max = -1L, na = character(),
  locale = default_locale(), progress = interactive())
```

Read Apache style log files

```
read_log(file, col_names = FALSE, col_types = NULL, skip = 0, n_max = -1, progress = interactive())
```

Read a file into a raw vector

```
read_file_raw(file)
```

Read each line into a raw vector

```
read_lines_raw(file, skip = 0, n_max = -1L,
  progress = interactive())
```

Data types

readr functions guess the types of each column and convert types when appropriate (but will NOT convert strings to factors automatically).

A message shows the type of each column in the result.

```
## Parsed with column specification:
## cols(
##   age = col_integer(),
##   sex = col_character(),
##   earn = col_double()
## )
```

earn is a double (numeric)

sex is a character

1. Use **problems()** to diagnose problems.

```
x <- read_csv("file.csv"); problems(x)
```

2. Use a **col_** function to guide parsing.

- **col_guess()** - the default
- **col_character()**
- **col_double()**, **col_euro_double()**
- **col_datetime()** (format = "") Also
- **col_date()** (format = ""), **col_time()** (format = "")
- **col_factor()** (levels, ordered = FALSE)
- **col_integer()**
- **col_logical()**
- **col_number()**, **col_numeric()**
- **col_skip()**

```
x <- read_csv("file.csv", col_types = cols(
  A = col_double(),
  B = col_logical(),
  C = col_factor()))
```

3. Else, read in as character vectors then parse with a **parse_** function.

- **parse_guess()**
 - **parse_character()**
 - **parse_datetime()** Also **parse_date()** and **parse_time()**
 - **parse_double()**
 - **parse_factor()**
 - **parse_integer()**
 - **parse_logical()**
 - **parse_number()**
- ```
x$A <- parse_number(x$A)
```



# Introducing `read_csv`

- Many options can be specified in `read_csv` (although usually don't need to be)

```
read_csv(file,
 col_names = TRUE,
 col_types = NULL,
 locale = default_locale(),
 na = c("", "NA"),
 quoted_na = TRUE,
 quote = "\"",
 comment = "",
 trim_ws = TRUE,
 skip = 0,
 n_max = Inf,
 guess_max = min(1000, n_max),
 progress = show_progress(),
 skip_empty_rows = TRUE)
```

```
bccdc_reportable_disease_age_sex - Notepad
File Edit Format View Help
Disease,Report Year,Sex,Age Group,Count
AIDS,2003,F,<1,0
AIDS,2003,F,1-4,0
AIDS,2003,F,10-14,0
AIDS,2003,F,15-19,1
AIDS,2003,F,20-24,1
AIDS,2003,F,25-29,1
AIDS,2003,F,30-39,4
AIDS,2003,F,40-59,13
AIDS,2003,F,5-9,0
AIDS,2003,F,60+,2
AIDS,2003,F,UNK,0
AIDS,2003,M,<1,0
AIDS,2003,M,1-4,0
AIDS,2003,M,10-14,0
AIDS,2003,M,15-19,0
AIDS,2003,M,20-24,1
AIDS,2003,M,25-29,8
AIDS,2003,M,30-39,53
AIDS,2003,M,40-59,61
AIDS,2003,M,5-9,0
AIDS,2003,M,60+,8
AIDS,2003,M,UNK,0
AIDS,2003,T,<1,0
AIDS,2003,T,1-4,0
```

Let's use the BCCDC reportable disease dashboard data by age and sex as an example

# Importing data with `read_csv`

- Most of the time, you only need to specify the file location and `read_csv` will helpfully do the rest


```
> library(tidyverse)
> indat <- "C:/Users/Michael/Google Drive/presentations/data import/data"
> infile <- "bccdc_reportable_disease_age_sex.txt"
>
> example.data <- read_csv(paste(indat, infile, sep = "/"))
Parsed with column specification:
cols(
 Disease = col_character(),
 `Report Year` = col_double(),
 Sex = col_character(),
 `Age Group` = col_character(),
 Count = col_double()
)
> |
```



# Importing data with `read_csv`

- Most of the time, you only need to specify the file location and `read_csv` will helpfully do the rest

```
> library(tidyverse)
> indat <- "C:/Users/Michael/Google Drive/presentations/data import/data"
> infile <- "bccdc_reportable_disease_age_sex.txt"
>
> example.data <- read_csv(paste(indat, infile, sep = "/"))
Parsed with column specification:
cols(
 Disease = col_character(),
 `Report Year` = col_double(),
 Sex = col_character(),
 `Age Group` = col_character(),
 Count = col_double()
)
> |
```



Note that `readr` has guessed the appropriate column formats, but tells you explicitly in case any are incorrect

# Importing data with `read_csv`

- What did `read_csv` do?

```
> example.data
A tibble: 19,411 x 5
 Disease `Report Year` Sex `Age Group` Count
 <chr> <dbl> <chr> <chr> <dbl>
1 AIDS 2003 F <1 0
2 AIDS 2003 F 1-4 0
3 AIDS 2003 F 10-14 0
4 AIDS 2003 F 15-19 1
5 AIDS 2003 F 20-24 1
6 AIDS 2003 F 25-29 1
7 AIDS 2003 F 30-39 4
8 AIDS 2003 F 40-59 13
9 AIDS 2003 F 5-9 0
10 AIDS 2003 F 60+ 2
... with 19,401 more rows
> |
```

# Importing data with `read_csv`

- What did `read_csv` do?

Imports data as a 'tibble' (user friendly version of a data frame)

```
> example.data
A tibble: 19,411 x 5
 Disease `Report Year` Sex `Age Group` Count
 <chr> <dbl> <chr> <chr> <dbl>
1 AIDS 2003 F <1 0
2 AIDS 2003 F 1-4 0
3 AIDS 2003 F 10-14 0
4 AIDS 2003 F 15-19 1
5 AIDS 2003 F 20-24 1
6 AIDS 2003 F 25-29 1
7 AIDS 2003 F 30-39 4
8 AIDS 2003 F 40-59 13
9 AIDS 2003 F 5-9 0
10 AIDS 2003 F 60+ 2
... with 19,401 more rows
> |
```

Preserves column names

# What about `read.csv`?

- Base R uses data import function `read.csv`

```
> example.data.old <- read.csv(paste(indat, infile, sep = "/"))
>
> example.data.old
```

|    | Disease | Report.Year | Sex | Age.Group | Count |
|----|---------|-------------|-----|-----------|-------|
| 1  | AIDS    | 2003        | F   | <1        | 0     |
| 2  | AIDS    | 2003        | F   | 1-4       | 0     |
| 3  | AIDS    | 2003        | F   | 10-14     | 0     |
| 4  | AIDS    | 2003        | F   | 15-19     | 1     |
| 5  | AIDS    | 2003        | F   | 20-24     | 1     |
| 6  | AIDS    | 2003        | F   | 25-29     | 1     |
| 7  | AIDS    | 2003        | F   | 30-39     | 4     |
| 8  | AIDS    | 2003        | F   | 40-59     | 13    |
| 9  | AIDS    | 2003        | F   | 5-9       | 0     |
| 10 | AIDS    | 2003        | F   | 60+       | 2     |
| 11 | AIDS    | 2003        | F   | UNK       | 0     |

# What about `read.csv`?

- What did `read.csv` do?

Imports data as a 'data frame'

```
> example.data.old <- read.csv(paste(indat, infile, sep = "/"))
```

```
>
```

```
> example.data.old
```

|    | Disease | Report.Year | Sex | Age.Group | Count |
|----|---------|-------------|-----|-----------|-------|
| 1  | AIDS    | 2003        | F   | <1        | 0     |
| 2  | AIDS    | 2003        | F   | 1-4       | 0     |
| 3  | AIDS    | 2003        | F   | 10-14     | 0     |
| 4  | AIDS    | 2003        | F   | 15-19     | 1     |
| 5  | AIDS    | 2003        | F   | 20-24     | 1     |
| 6  | AIDS    | 2003        | F   | 25-29     | 1     |
| 7  | AIDS    | 2003        | F   | 30-39     | 4     |
| 8  | AIDS    | 2003        | F   | 40-59     | 13    |
| 9  | AIDS    | 2003        | F   | 5-9       | 0     |
| 10 | AIDS    | 2003        | F   | 60+       | 2     |
| 11 | AIDS    | 2003        | F   | UNK       | 0     |

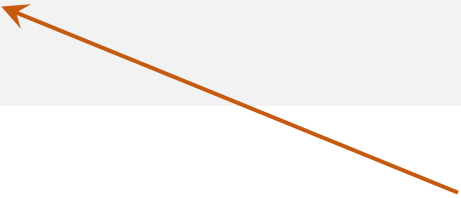
May change  
column names

Doesn't indicate,  
dataset size or  
column formats

# What about `read.csv`?

- What did `read.csv` do?

```
> glimpse(example.data.old)
Observations: 19,411
Variables: 5
$ Disease <fct> AIDS, AIDS, AIDS, AIDS, AIDS, AIDS, AIDS, AIDS, AIDS
$ Report.Year <int> 2003, 2003, 2003, 2003, 2003, 2003, 2003, 2003, 2003
$ Sex <fct> F, F, F, F, F, F, F, F, F, F, F, F, M, M, M, M, M, M
$ Age.Group <fct> <1, 1-4, 10-14, 15-19, 20-24, 25-29, 30-39, 40-59, 5
$ Count <int> 0, 0, 0, 1, 1, 1, 4, 13, 0, 2, 0, 0, 0, 0, 0, 1, 8,
> |
```



Imported character variables as factors (which can cause problems later and should generally be avoided)

# What about `read.csv`?

- Compared to the base R function `read.csv`, `read_csv`
  - retains original column names
  - is much faster (up to 10x)
  - does not convert character strings to factors by default
  - automatically parses common date/time formats
  - has a progress bar if loading is slow

# A messier example

Here is the data file I promised you

note that I did not include the column names so these will need to be added manually

also I have put some comments throughout to help clarify the data elements

# here are the counts by HA

FHA,2015-01-01,1,5,3

IHA,2015-05-01,4,\*missing\*,9

NHA,2016-03-13,6,8,2

# here are the corresponding ages

FHA,2015-01-01,34,45,13

IHA,2015-05-01,56,N/A,23

NHA,2016-03-13,56,23,63

What if we receive the data as a text file or spreadsheet with unneeded text and comments interspersed with the data?



# A messier example

|    | A                                                                                    | B          | C  | D         | E  | F | G | H |
|----|--------------------------------------------------------------------------------------|------------|----|-----------|----|---|---|---|
| 1  | Here is the data file I promised you                                                 |            |    |           |    |   |   |   |
| 2  | note that I did not include the column names so these will need to be added manually |            |    |           |    |   |   |   |
| 3  | also I have put some comments throughout to help clarify the data elements           |            |    |           |    |   |   |   |
| 4  |                                                                                      |            |    |           |    |   |   |   |
| 5  | # here are the counts by HA                                                          |            |    |           |    |   |   |   |
| 6  |                                                                                      |            |    |           |    |   |   |   |
| 7  | FHA                                                                                  | 2015-01-01 | 1  | 5         | 3  |   |   |   |
| 8  | IHA                                                                                  | 2015-05-01 | 4  | *missing* | 9  |   |   |   |
| 9  | NHA                                                                                  | 2016-03-13 | 6  | 8         | 2  |   |   |   |
| 10 |                                                                                      |            |    |           |    |   |   |   |
| 11 | # here are the corresponding ages                                                    |            |    |           |    |   |   |   |
| 12 |                                                                                      |            |    |           |    |   |   |   |
| 13 | FHA                                                                                  | 2015-01-01 | 34 | 45        | 13 |   |   |   |
| 14 | IHA                                                                                  | 2015-05-01 | 56 | N/A       | 23 |   |   |   |
| 15 | NHA                                                                                  | 2016-03-13 | 56 | 23        | 63 |   |   |   |

What if we receive the data as a text file or spreadsheet with unneeded text and comments interspersed with the data?

# A messier example

- We can make use of optional settings in `read_csv`

```
read_csv(file,
 col_names = TRUE,
 col_types = NULL,
 locale = default_locale(),
 na = c("", "NA"),
 quoted_na = TRUE,
 quote = "\"",
 comment = "",
 trim_ws = TRUE,
 skip = 0,
 n_max = Inf,
 guess_max = min(1000, n_max),
 progress = show_progress(),
 skip_empty_rows = TRUE)
```

# A messier example

- We can make use of optional settings in `read_csv`

```
read_csv(file,
 col_names = TRUE, ← add column names
 col_types = NULL,
 locale = default_locale(),
 na = c("", "NA"), ← specify codes indicating missing data
 quoted_na = TRUE,
 quote = "\"",
 comment = "", ← specify codes indicating comments
 trim_ws = TRUE,
 skip = 0, ← number of rows to skip at beginning
 n_max = Inf,
 guess_max = min(1000, n_max),
 progress = show_progress(),
 skip_empty_rows = TRUE) ← drop rows with no data
```

# A messier example

```
> clean <- read_csv(messy,
+ col_names = c("HA", "Date", "Early", "Mid", "Late"),
+ skip = 6,
+ comment = "#",
+ na = c("*missing*", "N/A"))
>
> clean
A tibble: 6 x 5
 HA Date Early Mid Late
 <chr> <date> <dbl> <dbl> <dbl>
1 FHA 2015-01-01 1 5 3
2 IHA 2015-05-01 4 NA 9
3 NHA 2016-03-13 6 8 2
4 FHA 2015-01-01 34 45 13
5 IHA 2015-05-01 56 NA 23
6 NHA 2016-03-13 56 23 63
> |
```

# A messier example

```
> clean <- read_csv(messy,
+ col_names = c("HA", "Date", "Early", "Mid", "Late"),
+ skip = 6,
+ comment = "#",
+ na = c("*missing*", "N/A"))
>
> clean
A tibble: 6 x 5
 HA Date Early Mid Late
 <chr> <date> <dbl> <dbl> <dbl>
1 FHA 2015-01-01 1 5 3
2 IHA 2015-05-01 4 NA 9
3 NHA 2016-03-13 6 8 2
4 FHA 2015-01-01 34 45 13
5 IHA 2015-05-01 56 NA 23
6 NHA 2016-03-13 56 23 63
> |
```

Specify column names

Skip beginning rows and comments

Specify codes used for missing data

Also, by default `read_csv` automatically skipped empty rows

# Dealing with dates

- `read_csv` recognizes dates that are uniformly presented (e.g., all in the form year, month, day)

```
> uniform.dates <- c("Dates", "1809-02-11", "1809-02-11", "1809/02/11")
> read_csv(uniform.dates)
A tibble: 3 x 1
 Dates
 <date>
1 1809-02-11
2 1809-02-11
3 1809-02-11
> |
```

# Dealing with dates

- `read_csv` recognizes dates that are uniformly presented (e.g., all in the form year, month, day)

```
> uniform.dates <- c("Dates", "1809-02-11", "1809-02-11", "1809/02/11")
> read_csv(uniform.dates)
A tibble: 3 x 1
 Dates
 <date> ← Properly formatted as a date
1 1809-02-11
2 1809-02-11
3 1809-02-11
> |
```

# Dealing with dates

- When dates are not uniformly presented, `read_csv` simply formats as character (i.e., preserves the original data) and let's you do the clean-up

```
> mixed.dates <- c("Dates", "1809-02-11", "1809/2/11", "11-02-1908")
> read_csv(mixed.dates)
A tibble: 3 x 1
 Dates
 <chr> ← when in doubt, formatted as a character
1 1809-02-11
2 1809/2/11
3 11-02-1908
> |
```



# Dealing with dates

- Fortunately, with tidy data tools like the `lubridate` package, working with dates is easy

```
> library(lubridate)

Attaching package: 'lubridate'

The following object is masked from 'package:base':

 date

> raw.dates <- c("Dates", "1809-02-11", "1809-2-11", "1809/02/11",
+ "1809/2/11", "1809:02:11")
> mixed.dates <- read_csv(raw.dates)
>
> mixed.dates
A tibble: 5 x 1
 Dates
 <chr>
1 1809-02-11
2 1809-2-11
3 1809/02/11
4 1809/2/11
5 1809:02:11
> |
```

# Dealing with dates

- We only need specify the ordering of the date components (e.g., `ymd` indicates year, month, day order), `lubridate` does the rest

```
> clean.dates <- mixed.dates %>%
+ mutate(Dates = ymd(Dates))
>
> clean.dates
A tibble: 5 x 1
 Dates
 <date> ← Now properly formatted as a date
1 1809-02-11
2 1809-02-11
3 1809-02-11
4 1809-02-11
5 1809-02-11
> |
```